

In [1]:

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 from sklearn.model_selection import train_test_split
7
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.neighbors import KNeighborsClassifier
10 from sklearn.svm import SVC
11 from sklearn.tree import DecisionTreeClassifier
12 from sklearn.ensemble import RandomForestClassifier
13 from sklearn.ensemble import GradientBoostingClassifier
14 from sklearn.ensemble import AdaBoostClassifier
15
16 from sklearn.metrics import classification_report, accuracy_score
17
18
19 import warnings
20 warnings.filterwarnings("ignore")

```

In [2]:

```

1 df = pd.read_csv("abalone.data")
2 df.head()

```

Out[2]:

	M	0.455	0.365	0.095	0.514	0.2245	0.101	0.15	15
0	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
1	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
2	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
3	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7
4	I	0.425	0.300	0.095	0.3515	0.1410	0.0775	0.120	8

In [3]:

```

1 cols = ["Sex", "Length", "Diameter", "Height", "WholeWeight", "ShuckedWeight", "VisceraWeight"]

```

In [4]:

```
1 df = pd.read_csv("abalone.data", header=None, names = cols)
2 df.head()
```

Out[4]:

	Sex	Length	Diameter	Height	WholeWeight	ShuckedWeight	VisceraWeight	ShellWeight	F
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	

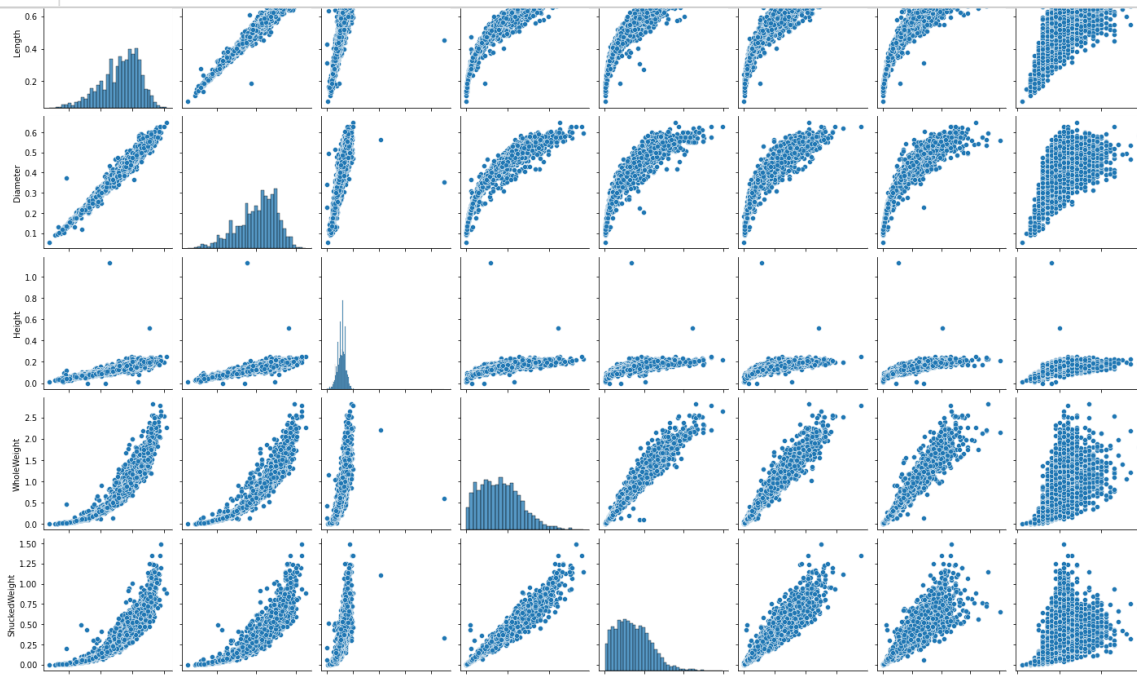
In [5]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   WholeWeight      4177 non-null   float64
5   ShuckedWeight    4177 non-null   float64
6   VisceraWeight     4177 non-null   float64
7   ShellWeight      4177 non-null   float64
8   Rings            4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

In [6]:

```
1 sns.pairplot(df)
```



In [7]:

```
1 x = df.iloc[:,1:].values
2 x
```

Out[7]:

```
array([[ 0.455 ,  0.365 ,  0.095 , ...,  0.101 ,  0.15 , 15.   ],
       [ 0.35  ,  0.265 ,  0.09  , ...,  0.0485,  0.07 ,  7.   ],
       [ 0.53  ,  0.42  ,  0.135 , ...,  0.1415,  0.21 ,  9.   ],
       ...,
       [ 0.6   ,  0.475 ,  0.205 , ...,  0.2875,  0.308 ,  9.   ],
       [ 0.625 ,  0.485 ,  0.15  , ...,  0.261 ,  0.296 , 10.   ],
       [ 0.71  ,  0.555 ,  0.195 , ...,  0.3765,  0.495 , 12.   ]])
```

In [8]:

```
1 from sklearn.cluster import KMeans
```

In [9]:

```
1 wcss = []
2
3 for i in range(1,11):
4     kmeans = KMeans(n_clusters=i, random_state=1)
5     kmeans.fit(x)
6     wcss.append(kmeans.inertia_)
```

In [10]:


```
1 wcss
```

Out[10]:

```
[44860.378975952,  
18861.998927978846,  
9889.828962494752,  
6340.081927459619,  
4384.82991662735,  
3438.389737637842,  
2706.8677666982144,  
2331.3658828700895,  
2021.483195630447,  
1717.8814779488623]
```

In [11]:

```
1 plt.plot(range(1,11), wcss)  
2 plt.title("The Elbow Method")  
3 plt.xlabel("Number of cluster")  
4 plt.ylabel("Wcss")  
5 plt.grid(True)  
6 plt.show()
```



In [12]:

```
1 kmeans = KMeans(n_clusters=3, random_state=1)  
2 ykmeans = kmeans.fit_predict(x)
```

In [13]:

```
1 ykmeans
```

Out[13]:

```
array([1, 2, 0, ..., 0, 0, 0])
```

In [14]:


```
1 kmeans.cluster_centers_
```

Out[14]:

```
array([[ 0.57364035,  0.44874342,  0.15363377,  1.00693048,  0.44330066,
         0.22172434,  0.28465482, 10.4127193 ],
       [ 0.58873469,  0.46642857,  0.16835714,  1.13776327,  0.4316051 ,
         0.23970306,  0.36281939, 16.46122449],
       [ 0.42099147,  0.32127576,  0.10659559,  0.4323742 ,  0.198199 ,
         0.09335714,  0.12139446,  6.88415068]])
```

In [15]:

```
1 plt.scatter(x[ykmeans==0,0], x[ykmeans==0,1], s=100, c="red", label="cluster1")
2 plt.scatter(x[ykmeans==1,0], x[ykmeans==1,1], s=100, c="blue", label="cluster2")
3 plt.scatter(x[ykmeans==2,0], x[ykmeans==2,1], s=100, c="green", label="cluster3")
4
5 plt.scatter(kmeans.cluster_centers_[ :,0], kmeans.cluster_centers_[ :,1], s=300, c="yellow")
6
7 plt.title("Clusters of Abalone ")
8 plt.xlabel("ShellWeight")
9 plt.ylabel("Rings")
10 plt.legend()
11 plt.show()
```



In [16]:

```
1 df["Target"] = ykmeans
```

In [17]:

1 df.head()

Out[17]:

	Sex	Length	Diameter	Height	WholeWeight	ShuckedWeight	VisceraWeight	ShellWeight	F
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	

In [18]:

1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   WholeWeight      4177 non-null   float64
5   ShuckedWeight    4177 non-null   float64
6   VisceraWeight    4177 non-null   float64
7   ShellWeight      4177 non-null   float64
8   Rings            4177 non-null   int64
9   Target           4177 non-null   int32
dtypes: float64(7), int32(1), int64(1), object(1)
memory usage: 310.1+ KB
```

In [19]:

1 df.isna().sum()

Out[19]:

```
Sex              0
Length           0
Diameter         0
Height           0
WholeWeight      0
ShuckedWeight    0
VisceraWeight    0
ShellWeight      0
Rings            0
Target           0
dtype: int64
```

In [20]:

```
1 sns.pairplot(df, hue="Target")
```

Out[20]:

<seaborn.axisgrid.PairGrid at 0xf36c455fd0>



In [21]:

```
1 x = df.iloc[:,1:].values  
2 y = df.iloc[:, -1].values
```

In [22]:

```
1 x
```

Out[22]:

```
array([[ 0.455,  0.365,  0.095, ...,  0.15 , 15.   ,  1.   ],
       [ 0.35 ,  0.265,  0.09 , ...,  0.07 ,  7.    ,  2.   ],
       [ 0.53 ,  0.42 ,  0.135, ...,  0.21 ,  9.    ,  0.   ],
       ...,
       [ 0.6   ,  0.475,  0.205, ...,  0.308,  9.    ,  0.   ],
       [ 0.625,  0.485,  0.15 , ...,  0.296, 10.   ,  0.   ],
       [ 0.71 ,  0.555,  0.195, ...,  0.495, 12.   ,  0.   ]])
```

In [23]:

```
1 y
```

Out[23]:

```
array([1, 2, 0, ..., 0, 0, 0])
```

In [24]:

```
1 xtrain,xtest,ytrain,ytest= train_test_split(x,y, test_size=0.3, random_state=1, stratif
```

In [25]:

```
1 def mymodel(model):
2     model.fit(xtrain,ytrain)
3     ypred = model.predict(xtest)
4     print(classification_report(ytest,ypred))
```

In [26]:

```
1 logreg = LogisticRegression()
2 knn = KNeighborsClassifier()
3 dt = DecisionTreeClassifier()
4 svm = SVC()
5 rf = RandomForestClassifier()
```


In [27]:

```

1 models = []
2 models.append(("KNN", KNeighborsClassifier()))
3 models.append(("Logreg", LogisticRegression()))
4 models.append(("SVM", SVC()))
5 models.append(("DT", DecisionTreeClassifier()))
6 models.append(("RF", RandomForestClassifier()))
7 models.append(("Ada", AdaBoostClassifier(n_estimators=100)))
8 models.append(("gbc", GradientBoostingClassifier(n_estimators=100)))
9 accuracy=[]
10
11 for name, model in models:
12     model.fit(xtrain, ytrain)
13     ypred = model.predict(xtest)
14     ac = accuracy_score(ytest, ypred)
15     accuracy.append(round(ac*100))
16
17     print(name)
18     print(classification_report(ytest, ypred))
19     print()
20
21
22 print(f"Avg. Ensemble Accuracy-: {np.array(accuracy).mean()} %")

```

KNN -:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	685
1	1.00	1.00	1.00	147
2	1.00	1.00	1.00	422
accuracy			1.00	1254
macro avg	1.00	1.00	1.00	1254
weighted avg	1.00	1.00	1.00	1254

Logreg -:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	685
1	1.00	1.00	1.00	147
2	1.00	1.00	1.00	422
accuracy			1.00	1254
macro avg	1.00	1.00	1.00	1254
weighted avg	1.00	1.00	1.00	1254

SVM -:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	685
1	1.00	1.00	1.00	147
2	1.00	1.00	1.00	422
accuracy			1.00	1254
macro avg	1.00	1.00	1.00	1254
weighted avg	1.00	1.00	1.00	1254

```
DT      -:
          precision    recall  f1-score   support

     0       1.00        1.00        1.00        685
     1       1.00        1.00        1.00        147
     2       1.00        1.00        1.00        422

 accuracy          1.00          1.00          1.00        1254
 macro avg          1.00          1.00          1.00        1254
weighted avg          1.00          1.00          1.00        1254
```

```
RF      -:
          precision    recall  f1-score   support

     0       1.00        1.00        1.00        685
     1       1.00        1.00        1.00        147
     2       1.00        1.00        1.00        422

 accuracy          1.00          1.00          1.00        1254
 macro avg          1.00          1.00          1.00        1254
weighted avg          1.00          1.00          1.00        1254
```

```
Ada     -:
          precision    recall  f1-score   support

     0       1.00        1.00        1.00        685
     1       1.00        1.00        1.00        147
     2       1.00        1.00        1.00        422

 accuracy          1.00          1.00          1.00        1254
 macro avg          1.00          1.00          1.00        1254
weighted avg          1.00          1.00          1.00        1254
```

```
gbc     -:
          precision    recall  f1-score   support

     0       1.00        1.00        1.00        685
     1       1.00        1.00        1.00        147
     2       1.00        1.00        1.00        422

 accuracy          1.00          1.00          1.00        1254
 macro avg          1.00          1.00          1.00        1254
weighted avg          1.00          1.00          1.00        1254
```

Avg. Ensemble Accuracy-: 100.0 %

Hyper Parameter Tuning

In [28]:

```

1 logregs = []
2
3 logregs.append(("LogregLin", LogisticRegression(solver = 'liblinear')))
4 logregs.append(("LogregLbf", LogisticRegression(solver = 'lbfgs')))
5 logregs.append(("LogregNcg", LogisticRegression(solver = 'newton-cg')))
6
7 for name, model in logregs:
8     print(name)
9     mymodel(model)
10    print("\n\n")
11

```

LogregLin		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

LogregLbf		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

LogregNcg		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

In [29]:

1	mymodel(logreg)				
	precision	recall	f1-score	support	
0	1.00	1.00	1.00	685	
1	1.00	1.00	1.00	147	
2	1.00	1.00	1.00	422	
accuracy			1.00	1254	
macro avg	1.00	1.00	1.00	1254	
weighted avg	1.00	1.00	1.00	1254	

In [30]:

```

1 svms = []
2 svms.append(("SVMlin", SVC(kernel = 'linear')))
3 svms.append(("SVMpol", SVC(kernel = 'poly')))
4 svms.append(("SVMrbf", SVC(kernel = 'rbf')))
5
6 for name, model in svms:
7     print(name)
8     mymodel(model)
9     print("\n\n\n")

```

SVMlin		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

SVMpol		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

SVMrbf		:-			
		precision	recall	f1-score	support
	0	1.00	1.00	1.00	685
	1	1.00	1.00	1.00	147
	2	1.00	1.00	1.00	422
	accuracy			1.00	1254
	macro avg	1.00	1.00	1.00	1254
	weighted avg	1.00	1.00	1.00	1254

In [31]:

```
1 mymodel(svm)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	685
1	1.00	1.00	1.00	147
2	1.00	1.00	1.00	422
accuracy			1.00	1254
macro avg	1.00	1.00	1.00	1254
weighted avg	1.00	1.00	1.00	1254

Cross Validation Score

In [32]:

```
1 from sklearn.model_selection import cross_val_score
```

In [33]:

```
1 for name, model in models:
2     print(name)
3     cvs = cross_val_score(model,x,y,cv=5,scoring='accuracy')
4     print(cvs.mean())
5     print("\n\n\n")
```

KNN - :
1.0

Logreg - :
1.0

SVM - :
1.0

DT - :
1.0

RF - :
1.0

Ada - :
1.0

gbc - :
1.0

Results

- This model gives 100% accuracy.
- Some HyperParameterTuning has been done even though not necessary
- the mean accuracy of the model is also 100%

In []:

1	
---	--